

miniprojmark

2023-05-08

```
# save csv file to project folder
fna.data <- "WisconsinCancer.csv"
#assign to a variable
wisc.df <- read.csv(fna.data, row.names = 1)
```

```
# remove expert diagnosis from data aka first column
wisc.data <- wisc.df[,-1]
#save first column for later
diagnosis <- factor(wisc.df[,1])
```

diagnosis

```
##      [1] M M M M M M M M M M M M M M M M M M M M M M M M M M M M M
##     [38] B M M M M M M M M M B B B B B M M B M M B B B B M B M M B B B
##     [75] B M B M M B B B M M B M M M B B B M B B M M B B B B M B B M B B
##    [112] B B B B B B M M M B M M B B B M M B M M M B M M B B M B B M B B M B
##    [149] B B B B B B B B M B B B B M M B M B B M M B B B B M B B M M M B M
##    [186] B M B B B M B B M M B M M M M B M M M B M B B M B M M M M B B M B B
##    [223] B M B B B B B M M B B B M B B M B B B B B M B B B B B M B M M M M M
##    [260] M M M M M M M B B B B B B M B M B B M B B M B M M B B B B B B B B B
##    [297] B M B B M B M B B B B B B B B B B B B B M B B B M B M B B B B M M B B
##    [334] B B M B M B M B B B M B B B B B B B M M M B B B B B B B B B B M M M
##    [371] M B M M B B B B B M B B B B B M B B B M B B M M B B B B B M B B B B B
##    [408] B M B B B B B M B B M B B B B B B B B B B M B M M B M B B B B M B B
##    [445] M B M B B M B M B B B B B B B B M M B B B B B B M B B B B B B M B
##    [482] B B B B B B M B M B B M B B B B B M M B M B M B B B B M B B M B M M
##    [519] B B B M B B B B B B B B B B B M B M M B B B B B B B B B B B B B B
##    [556] B B B B B B B M M M M M M B
## Levels: B M
```

Q1. How many observations are there? 569 observations

Q2. How many observations are malignant? 212 observations

```
grep("M", diagnosis)
```

```
##      [1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18
##     [19] 19 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 39 40
##     [37] 41 42 43 44 45 46 48 54 55 57 58 63 65 66 71 73 74 76
##     [55] 78 79 83 84 86 87 88 92 95 96 100 101 106 109 118 119 120 122
##     [73] 123 127 128 130 132 133 135 136 139 142 147 157 162 163 165 168 169 172
##     [91] 173 178 181 182 183 185 187 191 194 195 197 198 199 200 202 203 204 206
##    [109] 208 211 213 214 215 216 219 220 224 230 231 234 237 238 240 245 251 253
```

```
## [127] 254 255 256 257 258 259 260 261 262 263 264 265 266 273 275 278 281 283
## [145] 284 298 301 303 318 322 324 329 330 331 336 338 340 344 352 353 354 366
## [163] 367 369 370 371 373 374 380 386 390 393 394 401 409 415 418 431 433 434
## [181] 436 442 445 447 450 452 461 462 469 480 488 490 493 499 500 502 504 510
## [199] 513 515 517 518 522 534 536 537 563 564 565 566 567 568
```

Q3. How many variables/features are suffixed with `_mean`? 10 variables

```
#check column means and standard dev for scaling purposes
colMeans(wisc.data)
```

```
##          radius_mean          texture_mean          perimeter_mean
##      1.412729e+01      1.928965e+01      9.196903e+01
##          area_mean          smoothness_mean          compactness_mean
##      6.548891e+02      9.636028e-02      1.043410e-01
##      concavity_mean      concave.points_mean          symmetry_mean
##      8.879932e-02      4.891915e-02      1.811619e-01
## fractal_dimension_mean          radius_se          texture_se
##      6.279761e-02      4.051721e-01      1.216853e+00
##          perimeter_se          area_se          smoothness_se
##      2.866059e+00      4.033708e+01      7.040979e-03
##      compactness_se          concavity_se      concave.points_se
##      2.547814e-02      3.189372e-02      1.179614e-02
##          symmetry_se      fractal_dimension_se          radius_worst
##      2.054230e-02      3.794904e-03      1.626919e+01
##          texture_worst          perimeter_worst          area_worst
##      2.567722e+01      1.072612e+02      8.805831e+02
##      smoothness_worst      compactness_worst          concavity_worst
##      1.323686e-01      2.542650e-01      2.721885e-01
##      concave.points_worst          symmetry_worst      fractal_dimension_worst
##      1.146062e-01      2.900756e-01      8.394582e-02
```

```
apply(wisc.data,2,sd)
```

```
##          radius_mean          texture_mean          perimeter_mean
##      3.524049e+00      4.301036e+00      2.429898e+01
##          area_mean          smoothness_mean          compactness_mean
##      3.519141e+02      1.406413e-02      5.281276e-02
##      concavity_mean      concave.points_mean          symmetry_mean
##      7.971981e-02      3.880284e-02      2.741428e-02
## fractal_dimension_mean          radius_se          texture_se
##      7.060363e-03      2.773127e-01      5.516484e-01
##          perimeter_se          area_se          smoothness_se
##      2.021855e+00      4.549101e+01      3.002518e-03
##      compactness_se          concavity_se      concave.points_se
##      1.790818e-02      3.018606e-02      6.170285e-03
##          symmetry_se      fractal_dimension_se          radius_worst
##      8.266372e-03      2.646071e-03      4.833242e+00
##          texture_worst          perimeter_worst          area_worst
##      6.146258e+00      3.360254e+01      5.693570e+02
##      smoothness_worst      compactness_worst          concavity_worst
##      2.283243e-02      1.573365e-01      2.086243e-01
##      concave.points_worst          symmetry_worst      fractal_dimension_worst
##      6.573234e-02      6.186747e-02      1.806127e-02
```

```
#perform pca on wisc.data
wisc.pr <- prcomp(wisc.data, scale=TRUE)
apply
```

```
## function (X, MARGIN, FUN, ..., simplify = TRUE)
## {
##   FUN <- match.fun(FUN)
##   simplify <- isTRUE(simplify)
##   dl <- length(dim(X))
##   if (!dl)
##     stop("dim(X) must have a positive length")
##   if (is.object(X))
##     X <- if (dl == 2L)
##       as.matrix(X)
##     else as.array(X)
##   d <- dim(X)
##   dn <- dimnames(X)
##   ds <- seq_len(dl)
##   if (is.character(MARGIN)) {
##     if (is.null(dnn <- names(dn)))
##       stop("'X' must have named dimnames")
##     MARGIN <- match(MARGIN, dnn)
##     if (anyNA(MARGIN))
##       stop("not all elements of 'MARGIN' are names of dimensions")
##   }
##   d.call <- d[-MARGIN]
##   d.ans <- d[MARGIN]
##   if (anyNA(d.call) || anyNA(d.ans))
##     stop("'MARGIN' does not match dim(X)")
##   s.call <- ds[-MARGIN]
##   s.ans <- ds[MARGIN]
##   dn.call <- dn[-MARGIN]
##   dn.ans <- dn[MARGIN]
##   d2 <- prod(d.ans)
##   if (d2 == 0L) {
##     newX <- array(vector(typeof(X), 1L), dim = c(prod(d.call),
##       1L))
##     ans <- forceAndCall(1, FUN, if (length(d.call) < 2L) newX[,
##       1] else array(newX[, 1L], d.call, dn.call), ...)
##     return(if (is.null(ans)) ans else if (length(d.ans) <
##       2L) ans[1L][-1L] else array(ans, d.ans, dn.ans))
##   }
##   newX <- aperm(X, c(s.call, s.ans))
##   dim(newX) <- c(prod(d.call), d2)
##   ans <- vector("list", d2)
##   if (length(d.call) < 2L) {
##     if (length(dn.call))
##       dimnames(newX) <- c(dn.call, list(NULL))
##     for (i in 1L:d2) {
##       tmp <- forceAndCall(1, FUN, newX[, i], ...)
##       if (!is.null(tmp))
##         ans[[i]] <- tmp
##     }
##   }
## }
```

```

##     }
##     else for (i in 1L:d2) {
##         tmp <- forceAndCall(1, FUN, array(newX[, i], d.call,
##             dn.call), ...)
##         if (!is.null(tmp))
##             ans[[i]] <- tmp
##     }
##     ans.list <- !simplify || is.recursive(ans[[1L]])
##     l.ans <- length(ans[[1L]])
##     ans.names <- names(ans[[1L]])
##     if (!ans.list)
##         ans.list <- any(lengths(ans) != l.ans)
##     if (!ans.list && length(ans.names)) {
##         all.same <- vapply(ans, function(x) identical(names(x),
##             ans.names), NA)
##         if (!all(all.same))
##             ans.names <- NULL
##     }
##     len.a <- if (ans.list)
##         d2
##     else length(ans <- unlist(ans, recursive = FALSE))
##     if (length(MARGIN) == 1L && len.a == d2) {
##         names(ans) <- if (length(dn.ans[[1L]]))
##             dn.ans[[1L]]
##         ans
##     }
##     else if (len.a == d2)
##         array(ans, d.ans, dn.ans)
##     else if (len.a && len.a%%d2 == 0L) {
##         if (is.null(dn.ans))
##             dn.ans <- vector(mode = "list", length(d.ans))
##         dn1 <- list(ans.names)
##         if (length(dn.call) && !is.null(n1 <- names(dn <- dn.call[1])) &&
##             nzchar(n1) && length(ans.names) == length(dn[[1]]))
##             names(dn1) <- n1
##         dn.ans <- c(dn1, dn.ans)
##         array(ans, c(len.a%%d2, d.ans), if (!is.null(names(dn.ans)) ||
##             !all(vapply(dn.ans, is.null, NA)))
##             dn.ans)
##     }
##     else ans
## }
## <bytecode: 0x0000014b50a72dd8>
## <environment: namespace:base>

```

```
summary(wisc.pr)
```

```

## Importance of components:
##           PC1    PC2    PC3    PC4    PC5    PC6    PC7
## Standard deviation    3.6444 2.3857 1.67867 1.40735 1.28403 1.09880 0.82172
## Proportion of Variance 0.4427 0.1897 0.09393 0.06602 0.05496 0.04025 0.02251
## Cumulative Proportion 0.4427 0.6324 0.72636 0.79239 0.84734 0.88759 0.91010
##           PC8    PC9    PC10   PC11   PC12   PC13   PC14
## Standard deviation    0.69037 0.6457 0.59219 0.5421 0.51104 0.49128 0.39624

```


```
## Proportion of Variance 0.01589 0.0139 0.01169 0.0098 0.00871 0.00805 0.00523
## Cumulative Proportion 0.92598 0.9399 0.95157 0.9614 0.97007 0.97812 0.98335
## PC15 PC16 PC17 PC18 PC19 PC20 PC21
## Standard deviation 0.30681 0.28260 0.24372 0.22939 0.22244 0.17652 0.1731
## Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
## Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
## PC22 PC23 PC24 PC25 PC26 PC27 PC28
## Standard deviation 0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
## Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
## Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
## PC29 PC30
## Standard deviation 0.02736 0.01153
## Proportion of Variance 0.00002 0.00000
## Cumulative Proportion 1.00000 1.00000
```

Q4. What proportion of the original variance is captured by PC1? 0.4427 or 44.27%

Q5. How many PCs required to describe atleast 70% of the original variance? 3 PCs are required.


Q6. How many PCs are required to describe atleast 90% of the original variance? 7 PCs are required.

```
#visualizing the PCA model
biplot(wisc.pr)
```

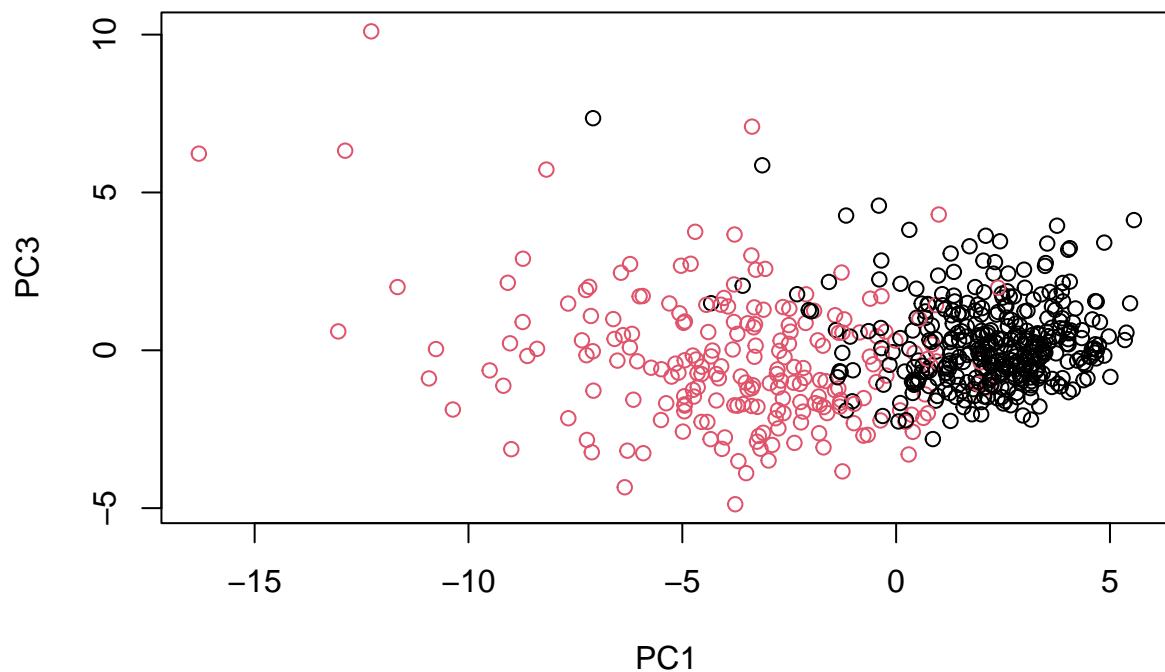


Q7. What stands out of this plot? I can't make much of this plot. It is very chaotic and full of overlapping data.

```
#scatter plot observations by component 1 and 2  
plot(wisc.pr$x[,1:2], col = diagnosis ,  
      xlab = "PC1", ylab = "PC2")
```



```
#scatter plot observations by component 1 and 3  
plot(wisc.pr$x[,1-3], col = diagnosis ,  
      xlab = "PC1", ylab = "PC3")
```

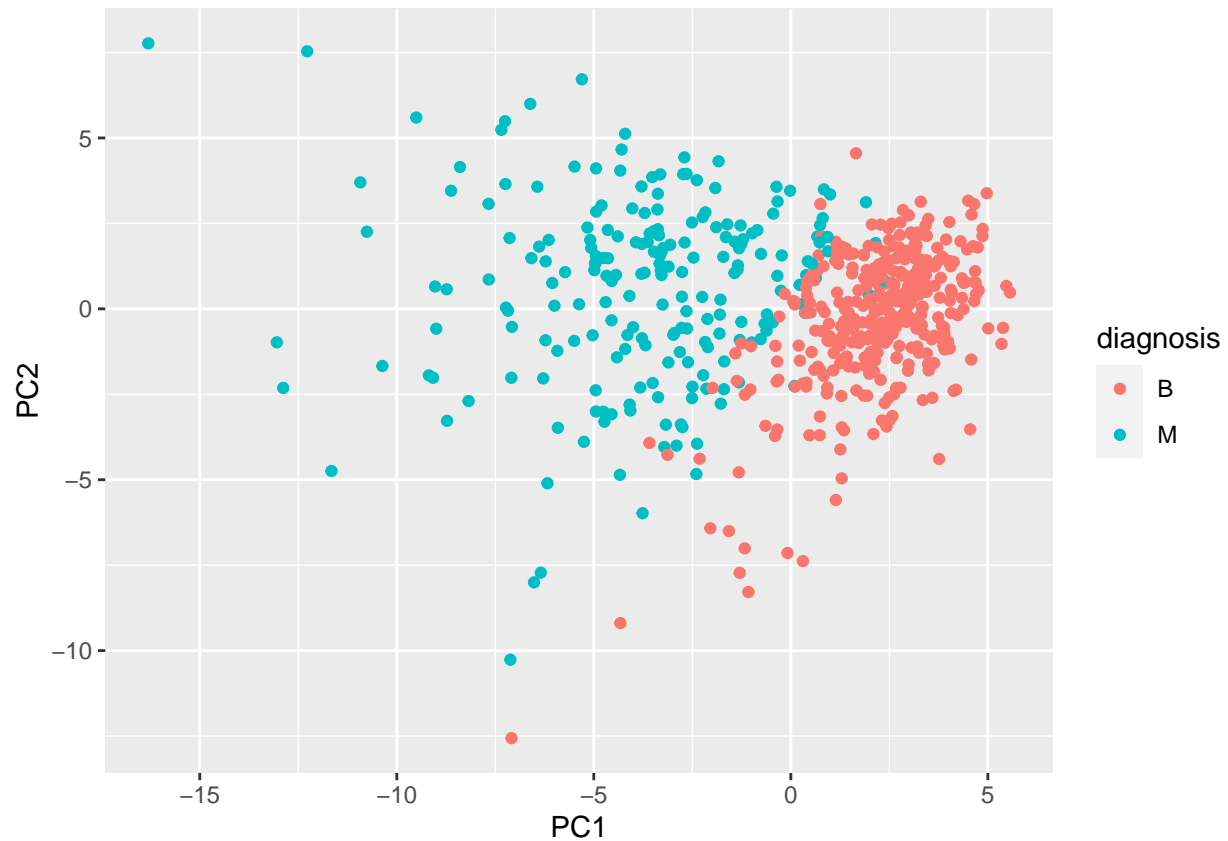


Q8. What do you notice about the plot for PC1/3? There seems to be a slightly clearer distinction between red and black as well as closer grouping.

```
# Create a data.frame for ggplot
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis

# Load the ggplot2 package
library(ggplot2)

# Make a scatter plot colored by diagnosis
ggplot(df) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```

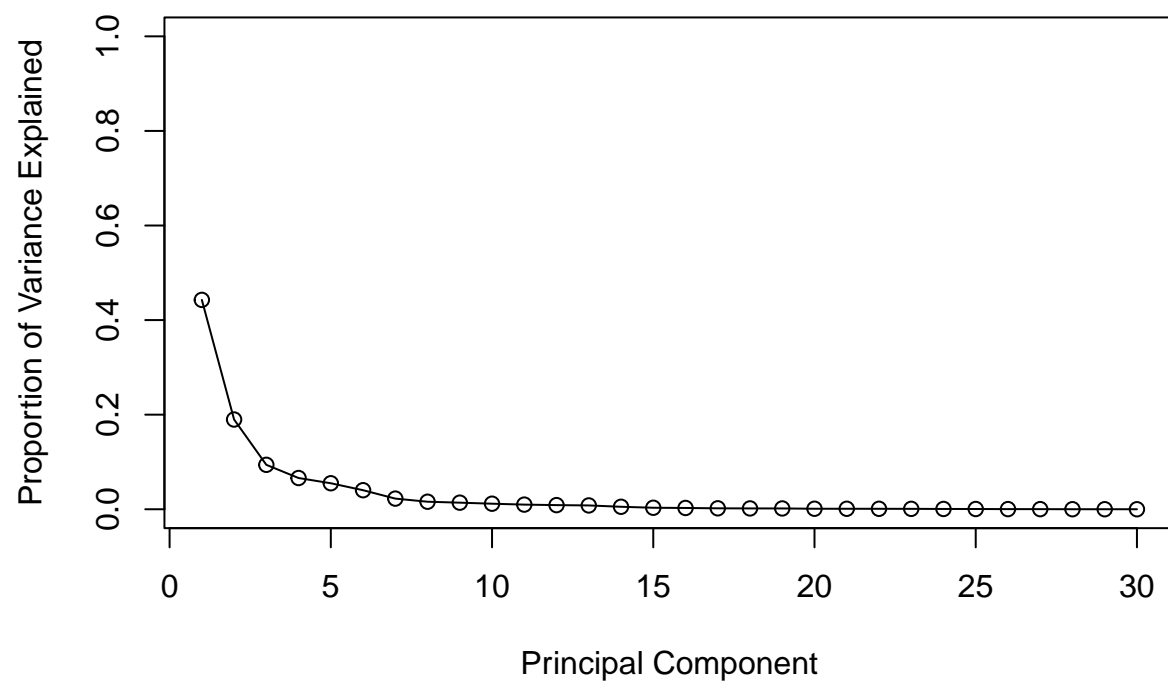


```
#calculate variance of each component
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

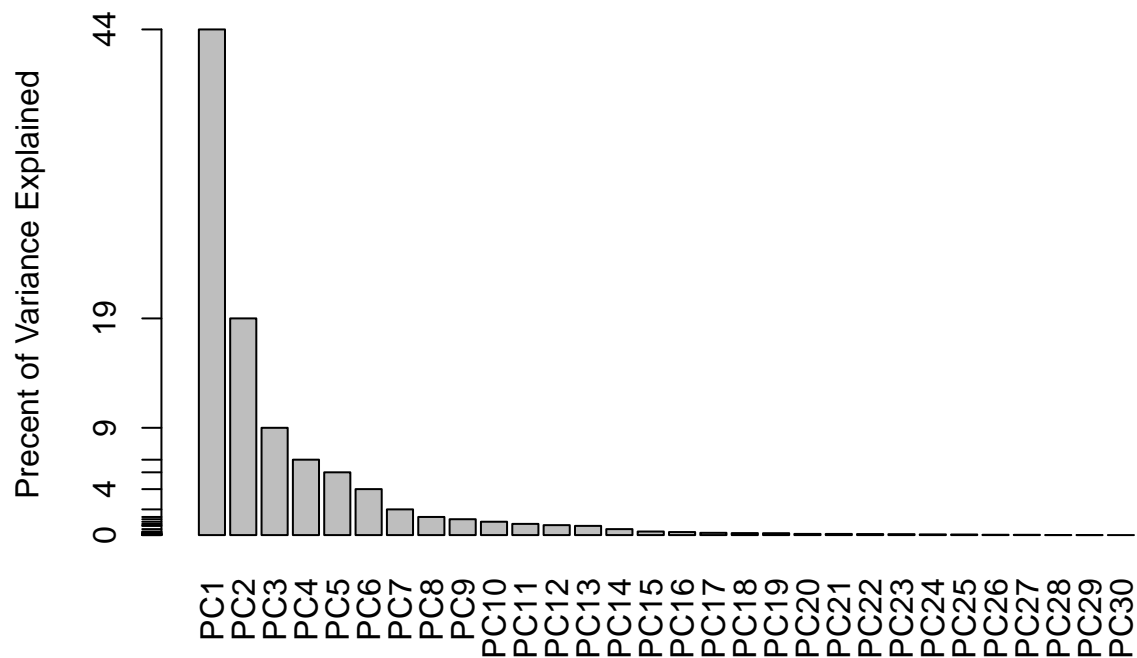
```
## [1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```

```
pve <- pr.var/sum(pr.var)

# plot variance
plot(pve, xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```

```
# Alternative scree plot of the same data, note data driven y-axis
barplot(pve, ylab = "Precent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```




Q9. What is the component of the loading vector for concave.points_mean for PC1? -0.26085376

Q10. Minimum number of PCs required to explain 80% of variance? 5 PCs

```
#scale wisc.data
data.scaled <- scale(wisc.data)
#calculate euclidean distance between all pairs of observations in scaled data
data.dist <- dist(data.scaled)
#create a heirarchal clustering model using complete linkage
wisc.hclust <- hclust(data.dist, method = "complete")
```

```
plot(wisc.hclust)
abline(wisc.hclust, col="red", lty=2)
```

Cluster Dendrogram



```
data.dist
hclust (*, "complete")
```

Q11. At what height is there 4 clusters? 20 (a horizontal line crosses 4 time at height 20)

```
#cutting so only 4 clusters
wisc.hclust.clusters <- cutree(wisc.hclust, k = 4)

table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##                   1 12 165
##                   2   2   5
##                   3 343  40
##                   4   0   2
```

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?


```
#cutting so only 4 clusters
wisc.hclust.clusterz <- cutree(wisc.hclust, k = 2)

table(wisc.hclust.clusterz, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusterz  B  M
##                   1 357 210
##                   2   0   2
```

Q12. Can you find a better cluster v diagnosis match by changing the # of clusters? Observing 2 clusters seems like a good match to observe. Less groups give us a clearer image of benign vs malignant.

```
wisc.hclustzz <- hclust(data.dist, method = "ward.D2")  
plot(wisc.hclustzz)
```




```
data.dist  
hclust (*, "ward.D2")
```

Q13. Favorite method? I think I like ward.D2 as well. Minimizing the variance creates such a clean and aesthetically pleasing picture to observe and understand.

```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[, 1:7]), method= "ward.D2")  
plot(wisc.pr.hclust)
```

Cluster Dendrogram



```
dist(wisc.pr$x[, 1:7])
hclust (*, "ward.D2")
```


```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
## grps
## 1 2
## 216 353
```


```
table(grps, diagnosis)
```

```
##      diagnosis
## grps    B    M
## 1    28 188
## 2   329  24
```

```
plot(wisc.pr$x[,1:2], col=grps)
```



```
plot(wisc.pr$x[,1:2], col=diagnosis)
```



```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[, 1:7]), method = "ward.D2")
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k = 2)
```

```
#compare diagnosis
table(wisc.pr.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.pr.hclust.clusters  B  M
##              1  28 188
##              2 329  24
```

Q15. How well does the newly created model with four clusters separate out the two diagnosis? I think it works really well. We are able to get a clearer picture of the diagnosis data and they are grouped together efficiently.

```
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##              1  12 165
##              2   2   5
##              3 343  40
##              4   0   2
```


Q16. How do the other models do in terms of separating diagnosis? H-cluster model does not work as well in separating diagnoses. The model we just made is more clear, has less groups to interpret. It is good to note that the models had similar outputs in terms of data.

Q17. Which analysis procedures resulted in the best specificity? Best sensitivity? I think the ward.D2 procedure had best specificity and sensitivity.

```
#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

```
##          PC1          PC2          PC3          PC4          PC5          PC6          PC7
## [1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
## [2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
##          PC8          PC9          PC10         PC11         PC12         PC13         PC14
## [1,] -0.2307350 0.1029569 -0.9272861 0.3411457  0.375921 0.1610764 1.187882
## [2,] -0.3307423 0.5281896 -0.4855301 0.7173233 -1.185917 0.5893856 0.303029
##          PC15         PC16         PC17         PC18         PC19         PC20
## [1,] 0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
## [2,] 0.1299153 0.1448061 -0.40509706 0.06565549 0.25591230 -0.4289500
##          PC21         PC22         PC23         PC24         PC25         PC26
## [1,] 0.1228233 0.09358453 0.08347651 0.1223396 0.02124121 0.078884581
## [2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
##          PC27         PC28         PC29         PC30
## [1,] 0.220199544 -0.02946023 -0.015620933 0.005269029
## [2,] -0.001134152 0.09638361 0.002795349 -0.019015820
```

```
plot(wisc.pr$x[,1:2])
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```

Q18. Which of the new patients should be prioritized for follow up based on results? Patient 2 needs to be followed up on for their results.